

Did They Hear Us?

A guide for civil society organisations to check whether their input shaped a governance process

Pedro dos Santos Maia

September 2026

Multilateral processes invite civil society to submit, and then produce their own documents. This guide shows how a coalition can check, using material that is already public, how much of what it wrote survived into what the institution wrote.

Why we wrote this

Your organisation spent months on a submission, and you do not know what happened to it.

This is the normal condition. Multilateral processes invite written input, publish the submissions, and then produce their own documents. Between the two there is no account of what was taken up, what was reworded and what was dropped. You are left with an impression, the secretariat has a different one, and nothing settles the matter.

Globethics' survey of 85 civil society organisations found technical capacity building to be the highest-ranked need, ahead of funding and coordination. This guide is one answer to that finding. It shows how a coalition can check, using documents that are already public, whether a process reflected what it collected.

We wrote it for the people who coordinate submissions rather than for researchers. You do not need to write code to use the first part.

What an audit can and cannot tell you

Read this before you start, so that you do not spend a week answering the wrong question.

It can tell you whether the themes in your submissions turn up in the institution's own documents, how much space each one takes compared with the space you gave it, and whether a demand survived intact or was kept in name while its content changed. It can tell you which of your asks left no trace. All of this rests on public documents, so anyone can check your working, including the people you are auditing.

It cannot tell you why. A theme may appear on an agenda because member states pushed it, because the secretariat always writes about it, or because it was planned before you wrote a word. The audit shows correspondence between two bodies of text. If you present that as proof of influence, you will lose the argument with the secretariat and your finding will be set aside.

It cannot tell you what was said in the room. Official documents describe how an agenda is framed. If the real work happened in corridors and side meetings, the audit will not see it.

So the claim stays narrow: this is how much of what we wrote appears in what they wrote. An institution that has publicly committed to participation still has to engage with that number, whatever caused it.

What this looked like at Geneva

We first ran the method on the UN's first Global Dialogue on AI Governance, held in July 2026. Three sets of public documents went in: 379 civil society submissions in English, 55 concept notes proposing side events, and the seven session notes the secretariat wrote for the official programme.

The themes came from what civil society wrote, rather than from a list decided in advance. We then measured each theme in all three collections on the same scale. The figures are shares of a collection's content, so they answer how much of this writing is about that, and not how many organisations mentioned it.

Theme	Civil society	Official programme
Inclusion and capacity-building	11.9%	25.8%
Data rights and corporate power	10.9%	3.5%
Legal accountability	6.1%	1.7%
Frontier AI risk	4.2%	1.0%

The pattern is consistent. The one theme that grew is the one asking for participation and capacity-building. Every theme asking for binding constraints on companies, governments or developers shrank. The words ban, moratorium and catastrophic appear in none of the seven official notes.

Two details matter more than the percentages. The sharpest filtering happened in the secretariat's own framing documents, and not at the side events, which tracked the submissions reasonably well. And one coalition had written, before the Dialogue, that its issue would be folded into general language unless somebody tracked it separately. It was folded in. They had predicted the outcome in their own submission, which is the kind of thing an audit can turn from a shared suspicion into something dated and checkable.

Choose your tier

Most guides of this kind assume you can run statistical software, and most organisations cannot. There are three levels here, and each one produces a real finding on its own. You do not need to reach the third to say something true.

Tier	What you need	Time	What you get
1. Word check	The documents, and a text editor	An afternoon	Whether your specific asks appear at all
2. Theme map	A spreadsheet and a careful reader	Two to three days	Which of your themes the agenda made room for
3. Full audit	Someone comfortable in R	A week	Shares per theme across every collection, comparable

Start at tier 1 even if you intend to reach tier 3. It takes an afternoon, it often produces the most quotable result, and it tells you whether the fuller work is worth doing.

One rule applies at every tier: **gather the documents before you look at any of them**. If you decide what counts as evidence after seeing which evidence helps you, the exercise is no longer an audit.

Tier 1 — the word check

The reasoning is simple. If you asked for a moratorium and the word moratorium appears nowhere in the agenda, the outcome document or the closing summary, that absence is a fact, and verifying it needs no method at all.

1. **Gather the institution's documents.** The agenda or programme, every session description, the outcome document, the chair's or co-chairs' summary. Save them as plain text or searchable PDF. Note how many there are and how long they are, since you will need to say so.
2. **Write down your defining terms before you search.** Take them from your own submission: the words you would not accept a paraphrase of. If you asked for binding rules, that is binding, enforceable, mandatory, obligation. If you asked for a halt, that is ban, moratorium, prohibition. Five to ten terms is enough. Writing them first stops you searching until something turns up.
3. **Search for whole words.** This matters more than it sounds. Searching for ban as a fragment will also match urban, and urban-rural divide appears in plenty of UN documents. Use your editor's whole-word option, or search for the term with a space on each side. If you report a hit you do not have, the whole audit loses its credibility.
4. **Record every result, including the ones that go against you.** A table of term, number

of documents containing it, number of occurrences. Terms that do appear belong in the finding as much as terms that do not.

5. **Read the hits in context.** A word can be present and still be doing nothing. Accountability in the phrase "accountability mechanisms could be explored" is some distance from a commitment. Quote the sentence rather than the count.

At Geneva this step produced the most repeated line of the whole analysis: ban, moratorium and catastrophic occurred in none of the seven official session notes. That took twenty minutes and no software.

Tier 2 — the theme map

Tier 1 tells you whether a word is there. Tier 2 tells you whether a theme is there, which is what matters when an institution has adopted your subject and dropped your ask.

1. **Build your list of themes from the submissions rather than from your own priorities.** Read thirty submissions chosen at random and write down what each is about, in your words. Stop adding new themes when three or four in a row produce nothing new. You will usually land on twelve to twenty. Doing this before you look at the agenda is the discipline the whole exercise rests on.
2. **Define each theme in one sentence**, specific enough that two people would code the same document the same way. "Corporate accountability" is too loose. "Asks for binding obligations on companies, as opposed to voluntary commitments" can be applied consistently.
3. **Code a sample.** Fifty submissions is enough to be useful. One row per document, one column per theme, mark present or absent. If two of you can code twenty of the same documents and compare, do it — the disagreements show which definitions are too loose, and fixing them early saves the finding later.
4. **Code the institution's documents the same way**, using the same definitions, without changing them. People tend to cut this step, and cutting it is why their audits fail to persuade.
5. **Compare proportions.** The agenda is shorter than the submissions, so raw counts will always fall. What you want is the share: of the themes present in civil society's writing, which ones did the agenda make room for? A theme in a quarter of submissions and in none of seven sessions is a finding. A theme that drops from forty mentions to three, on its own, is not.
6. **Classify what became of each theme.** Four outcomes cover almost everything:

Outcome	What it means
Amplified	More prominent in the agenda than in the submissions
Diluted	Present, but the specific ask replaced by something weaker
Engulfed	Absorbed into a broader item that changes what it is about
Excluded	No meaningful presence

Write down your reasoning for each classification as you go. When somebody disputes a call later, that note is all you will have.

Tier 3 — running the Echo Audit

What this adds over tier 2 is scale and comparability. Tier 2 codes a sample by hand; tier 3 covers every document and puts all collections on one scale, so that a share in the submissions and a share in the agenda can be compared directly. It also removes any suspicion that you picked the sample which suited you.

The code is published and runs from a clean copy. Install with `Rscript setup.R`, then run the worked Geneva example first. If your numbers match the ones recorded in the repository, your setup is correct and you can trust what it does with your own documents.

Prepare three spreadsheets, saved as CSV, one row per document: your community's submissions, the institution's own documents, and anything in between such as side-event proposals. Each needs an identifier column and one or more columns of text. If your form had several questions, keep them in separate columns, since the code joins them for you.

Phase 1, Voice finds the themes in the submissions. It does not use a list you supply, and that is deliberate: coding your community's words against categories you had chosen in advance would repeat the problem you are investigating.

```
source("R/voice.R")

voice <- run_echo_voice(

  csv_path  = "submissions.csv",

  text_cols = c("q1", "q2", "q3"),
```

```
id_col = "ID",

k_candidates = c(10, 15, 20, 25)

)
```

You now stop and read. The software fits several models with different numbers of themes and cannot tell you which is right. Look at the word lists for each. With too few themes, separate issues merge into one and share no distinguishing words. With too many, one issue splits across near-identical themes. Pick the version where the themes read as recognisable arguments, then read the documents the model says are most typical of each theme and name it from what they actually say.

This step cannot be automated, and it should not be. A number chosen by a statistic instead of a reading will produce themes that nobody recognises, including the people who wrote the submissions.

```
final <- finalize_topic_model(voice, k = 20)
```

Phase 2, Echo freezes that model and applies it, unchanged, to the institution's documents. The frozen model works like a ruler built from one set of texts and then carried to another, so the institution's text gets measured on your community's categories instead of its own.

```
source("R/echo.R")

official <- echo_project(voice, final, "sessions.csv",

  text_cols = "text_string", id_col = "ID",

  collection_name = "Official programme")

echo_compare(voice, final, list(official))
```

Check the **vocabulary coverage** that the output reports. It tells you what share of the institution's words your instrument recognises. High coverage, above roughly 85 per cent, means the comparison is sound. Low coverage means the two bodies of text have little language in common, and you will need to explain that before using the result.

Then classify each theme as amplified, diluted, engulfed or excluded, exactly as in tier 2. The software produces the shares, and you still have to make the judgement.

Three traps

Each of these has sunk a finding. The first two are mistakes you can make yourself; the third is an argument that will be made against you.

Repeat filings inflate a theme. Coalitions often file the same agreed text under several signatories' names. This is ordinary practice and nobody is hiding anything, but to a counting method one argument filed four times counts four times. At Geneva, a single coalition submission filed four times accounted for about a third of the theme it concerned.

The consequence is uncomfortable and worth sitting with. Coordinating thirteen signatories and filing repeatedly takes organisational capacity that most civil society organisations do not have. An uncorrected audit therefore weights the result toward the best-resourced submitters, and ends up reproducing the exclusion it was built to detect. Check for near-identical submissions before you measure anything, then decide openly whether to keep them, drop the duplicates, or treat the cluster as a single voice.

"An agenda has to be shorter." This is the first objection you will hear, and it sounds decisive. Four hundred submissions cannot fit into seven session notes, so of course things shrank.

It does not hold, and knowing why is worth more than any other paragraph in this guide. You are comparing shares rather than amounts. Within each collection the shares add up to the same total, so cutting everything equally changes nothing at all. A share only moves when attention was unequal. What you found is that some themes were cut far more than others.

House style. Institutions have habitual vocabulary. If a UN body writes about capacity-building in every document it ever produces, finding capacity-building in this one proves very little. The test is cheap: run your instrument over documents from an unrelated process by the same institution. If the amplified theme is just as prominent there, you are looking at house style.

There is a useful asymmetry here. House style can explain why something is present. It cannot explain why something is absent. Your findings about themes that shrank or disappeared hold whether or not you run this check.

Using what you find

An audit is only worth the days it costs if it changes a conversation. Four things make a finding hold up when it is contested.

Publish the working as well as the result. Say which documents you used, how many, when you downloaded them, and what you searched for. If you coded by hand, publish the

definitions. A finding people can check is much harder to dismiss than one they have to take on trust, and in our experience the checking rarely happens once it becomes possible.

Report what went against you. If three of your themes were taken up and one was dropped, say so. An audit that finds only what its authors hoped for reads as advocacy and gets treated as such. The Geneva analysis reported that one of its own headline findings was weaker than it first appeared, and the report was stronger for it.

Lead with absence. Of everything an audit produces, the most robust result is a defining term that appears nowhere. It needs no model, no sample and no statistics, and there is no counter-argument available except to change the document.

Go to the secretariat before you go public, if you intend to keep working with them. Show them the finding and ask whether they read it differently. Sometimes they have an explanation you did not have. When they do not, you have at least given them the chance, which counts for something the next time round.

One use gets overlooked. Run the audit on your own coalition's submissions before the next consultation opens. If your central ask is spread thinly across eleven documents in eleven different formulations, no instrument will find it concentrated, and neither will a secretariat.

Code, data and reuse

Everything described here is public, and the Geneva analysis can be reproduced from scratch.

- **Code and data:** github.com/pedrodsmaia/echo — the protocol, the R implementation of all three phases, and the Geneva pilot with its three corpora and the output you should get.
- **Permanent link:** doi.org/10.5281/zenodo.22643933 — archived and citable.

The code is MIT licensed and the protocol is CC BY 4.0, so you may use, adapt and republish both, including commercially, provided you credit the source. The Geneva corpora are public UN documents, redistributed so that the analysis can be checked.

The exact package versions we used are recorded in the repository. If your numbers do not match the published ones, a version has changed somewhere, and that is the first place to look before you start doubting your data.

The method extends the computational grounded theory framework of Laura K. Nelson (2020), *Sociological Methods & Research* 49(1), 3–42, which sets out the division of labour between machine pattern-detection and human reading that phases one and two follow. Please cite it

alongside this work.

If you run an audit on another process, the repository's issue tracker is the place to report what worked and what did not. We calibrated the thresholds in the protocol on two corpora, and they will need adjusting for very different material.